

People:

- **ID**: Customer's unique identifier
- **Year_Birth**: Customer's birth year
- **Education**: Customer's education level
- **Marital_Status**: Customer's marital status
- **Income**: Customer's yearly household income
- **Kidhome**: Number of children in customer's household
- **Teenhome**: Number of teenagers in customer's household
- **Dt_Customer**: Date of customer's enrollment with the company
- **Recency**: Number of days since customer's last purchase
- **Complain**: 1 if the customer complained in the last 2 years, 0 otherwise

Products:

- **MntWines**: Amount spent on wine in last 2 years
- **MntFruits**: Amount spent on fruits in last 2 years
- **MntMeatProducts**: Amount spent on meat in last 2 years
- **MntFishProducts**: Amount spent on fish in last 2 years
- **MntSweetProducts**: Amount spent on sweets in last 2 years
- **MntGoldProds**: Amount spent on gold in last 2 years

Promotion:

- **NumDealsPurchases**: Number of purchases made with a discount
- **AcceptedCmp1**: 1 if customer accepted the offer in the 1st campaign, 0 otherwise
- **AcceptedCmp2**: 1 if customer accepted the offer in the 2nd campaign, 0 otherwise
- **AcceptedCmp3**: 1 if customer accepted the offer in the 3rd campaign, 0 otherwise
- **AcceptedCmp4**: 1 if customer accepted the offer in the 4th campaign, 0 otherwise
- **AcceptedCmp5**: 1 if customer accepted the offer in the 5th campaign, 0 otherwise
- **Response**: 1 if customer accepted the offer in the last campaign, 0 otherwise

Place:

- **NumWebPurchases**: Number of purchases made through the company's website
- **NumCatalogPurchases**: Number of purchases made using a catalogue
- **NumStorePurchases**: Number of purchases made directly in stores
- **NumWebVisitsMonth**: Number of visits to company's website in the last month

[Back to Table of Contents](#)

```
# This Python 3 environment comes with many helpful analytics  
libraries installed  
# It is defined by the kaggle/python Docker image:  
https://github.com/kaggle/docker-python  
# For example, here's several helpful packages to load  
  
import numpy as np # linear algebra  
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)  
  
# Input data files are available in the read-only "../input/"  
directory  
# For example, running this (by clicking run or pressing Shift+Enter)  
will list all files under the input directory  
  
import os  
for dirname, _, filenames in os.walk('/kaggle/input'):  
    for filename in filenames:  
        print(os.path.join(dirname, filename))  
  
# You can write up to 20GB to the current directory (/kaggle/working/)  
that gets preserved as output when you create a version using "Save &  
Run All"  
# You can also write temporary files to /kaggle/temp/, but they won't  
be saved outside of the current session  
  
/kaggle/input/marketingdata/marketing_data.csv  
  
# handle table-like data and matrices  
import pandas as pd  
import numpy as np  
  
# visualisation  
import seaborn as sns  
import matplotlib.pyplot as plt  
import missingno as msno  
import plotly.express as px  
import plotly.graph_objects as go  
from plotly.subplots import make_subplots  
import plotly.figure_factory as ff  
  
# preprocessing  
from sklearn.preprocessing import StandardScaler  
  
# pca  
from sklearn.decomposition import PCA
```

```

# clustering
from yellowbrick.cluster import KElbowVisualizer
from sklearn.metrics import silhouette_score
from sklearn.cluster import KMeans, AgglomerativeClustering
from sklearn.cluster import DBSCAN

# evaluations
from sklearn.metrics import confusion_matrix

# ignore warnings
import warnings
warnings.filterwarnings('ignore')

# to display the total number columns present in the dataset
pd.set_option('display.max_columns', None)

/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146:
UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this
version of SciPy (detected version 1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and
<{np_maxversion}")

def wrangle(path):
    return pd.read_csv(path)

path = '/kaggle/input/marketingdata/marketing_data.csv'
df = wrangle(path)

```

[Back to Table of Contents](#)

```
df.head()
```

| | ID | Year_Birth | Education | Marital_Status | Income | Kidhome |
|---|-------|------------|------------|----------------|-------------|---------|
| 0 | 1826 | 1970 | Graduation | Divorced | \$84,835.00 | 0 |
| 1 | 1 | 1961 | Graduation | Single | \$57,091.00 | 0 |
| 2 | 10476 | 1958 | Graduation | Married | \$67,267.00 | 0 |
| 3 | 1386 | 1967 | Graduation | Together | \$32,474.00 | 1 |
| 4 | 5371 | 1989 | Graduation | Single | \$21,474.00 | 1 |

| | Teenhome | Dt_Customer | Recency | MntWines | MntFruits | MntMeatProducts |
|---|----------|-------------|---------|----------|-----------|-----------------|
| \ | | | | | | |

| | | | | | | |
|---|---|---------|---|-----|-----|-----|
| 0 | 0 | 6/16/14 | 0 | 189 | 104 | 379 |
| 1 | 0 | 6/15/14 | 0 | 464 | 5 | 64 |
| 2 | 1 | 5/13/14 | 0 | 134 | 11 | 59 |
| 3 | 1 | 5/11/14 | 0 | 10 | 0 | 1 |
| 4 | 0 | 4/8/14 | 0 | 6 | 16 | 24 |

| | MntFishProducts | MntSweetProducts | MntGoldProds | NumDealsPurchases |
|---|-----------------|------------------|--------------|-------------------|
| 0 | 111 | 189 | 218 | 1 |
| 1 | 7 | 0 | 37 | 1 |
| 2 | 15 | 2 | 30 | 1 |
| 3 | 0 | 0 | 0 | 1 |
| 4 | 11 | 0 | 34 | 2 |

| | NumWebPurchases | NumCatalogPurchases | NumStorePurchases |
|---|-----------------|---------------------|-------------------|
| 0 | 4 | 4 | 6 |
| 1 | | | |
| 1 | 7 | 3 | 7 |
| 5 | | | |
| 2 | 3 | 2 | 5 |
| 2 | | | |
| 3 | 1 | 0 | 2 |
| 7 | | | |
| 4 | 3 | 1 | 2 |
| 7 | | | |

| | AcceptedCmp3 AcceptedCmp2 \ | AcceptedCmp4 | AcceptedCmp5 | AcceptedCmp1 |
|---|--------------------------------|--------------|--------------|--------------|
| 0 | 0 | 0 | 0 | 0 |
| 0 | | | | |
| 1 | 0 | 0 | 0 | 0 |
| 1 | | | | |
| 2 | 0 | 0 | 0 | 0 |
| 0 | | | | |
| 3 | 0 | 0 | 0 | 0 |
| 0 | | | | |
| 4 | 1 | 0 | 0 | 0 |
| 0 | | | | |

| Response | Complain | Country |
|----------|----------|---------|
|----------|----------|---------|

| | | | |
|---|---|---|-----|
| 0 | 1 | 0 | SP |
| 1 | 1 | 0 | CA |
| 2 | 0 | 0 | US |
| 3 | 0 | 0 | AUS |
| 4 | 1 | 0 | SP |

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2240 entries, 0 to 2239
```

```
Data columns (total 28 columns):
```

| # | Column | Non-Null Count | Dtype |
|----|---------------------|----------------|--------|
| 0 | ID | 2240 non-null | int64 |
| 1 | Year_Birth | 2240 non-null | int64 |
| 2 | Education | 2240 non-null | object |
| 3 | Marital_Status | 2240 non-null | object |
| 4 | Income | 2216 non-null | object |
| 5 | Kidhome | 2240 non-null | int64 |
| 6 | Teenhome | 2240 non-null | int64 |
| 7 | Dt_Customer | 2240 non-null | object |
| 8 | Recency | 2240 non-null | int64 |
| 9 | MntWines | 2240 non-null | int64 |
| 10 | MntFruits | 2240 non-null | int64 |
| 11 | MntMeatProducts | 2240 non-null | int64 |
| 12 | MntFishProducts | 2240 non-null | int64 |
| 13 | MntSweetProducts | 2240 non-null | int64 |
| 14 | MntGoldProds | 2240 non-null | int64 |
| 15 | NumDealsPurchases | 2240 non-null | int64 |
| 16 | NumWebPurchases | 2240 non-null | int64 |
| 17 | NumCatalogPurchases | 2240 non-null | int64 |
| 18 | NumStorePurchases | 2240 non-null | int64 |
| 19 | NumWebVisitsMonth | 2240 non-null | int64 |
| 20 | AcceptedCmp3 | 2240 non-null | int64 |
| 21 | AcceptedCmp4 | 2240 non-null | int64 |
| 22 | AcceptedCmp5 | 2240 non-null | int64 |
| 23 | AcceptedCmp1 | 2240 non-null | int64 |
| 24 | AcceptedCmp2 | 2240 non-null | int64 |
| 25 | Response | 2240 non-null | int64 |
| 26 | Complain | 2240 non-null | int64 |
| 27 | Country | 2240 non-null | object |

```
dtypes: int64(23), object(5)
```

```
memory usage: 490.1+ KB
```

- We see that Income column has some missing values

```
df.describe().T
```

| | count | mean | std | min | 25% |
|---|-------|------|-----|-----|-----|
| \ | | | | | |

| | | | | | |
|---------------------|--------|-------------|-------------|--------|---------|
| ID | 2240.0 | 5592.159821 | 3246.662198 | 0.0 | 2828.25 |
| Year_Birth | 2240.0 | 1968.805804 | 11.984069 | 1893.0 | 1959.00 |
| Kidhome | 2240.0 | 0.444196 | 0.538398 | 0.0 | 0.00 |
| Teenhome | 2240.0 | 0.506250 | 0.544538 | 0.0 | 0.00 |
| Recency | 2240.0 | 49.109375 | 28.962453 | 0.0 | 24.00 |
| MntWines | 2240.0 | 303.935714 | 336.597393 | 0.0 | 23.75 |
| MntFruits | 2240.0 | 26.302232 | 39.773434 | 0.0 | 1.00 |
| MntMeatProducts | 2240.0 | 166.950000 | 225.715373 | 0.0 | 16.00 |
| MntFishProducts | 2240.0 | 37.525446 | 54.628979 | 0.0 | 3.00 |
| MntSweetProducts | 2240.0 | 27.062946 | 41.280498 | 0.0 | 1.00 |
| MntGoldProds | 2240.0 | 44.021875 | 52.167439 | 0.0 | 9.00 |
| NumDealsPurchases | 2240.0 | 2.325000 | 1.932238 | 0.0 | 1.00 |
| NumWebPurchases | 2240.0 | 4.084821 | 2.778714 | 0.0 | 2.00 |
| NumCatalogPurchases | 2240.0 | 2.662054 | 2.923101 | 0.0 | 0.00 |
| NumStorePurchases | 2240.0 | 5.790179 | 3.250958 | 0.0 | 3.00 |
| NumWebVisitsMonth | 2240.0 | 5.316518 | 2.426645 | 0.0 | 3.00 |
| AcceptedCmp3 | 2240.0 | 0.072768 | 0.259813 | 0.0 | 0.00 |
| AcceptedCmp4 | 2240.0 | 0.074554 | 0.262728 | 0.0 | 0.00 |
| AcceptedCmp5 | 2240.0 | 0.072768 | 0.259813 | 0.0 | 0.00 |
| AcceptedCmp1 | 2240.0 | 0.064286 | 0.245316 | 0.0 | 0.00 |
| AcceptedCmp2 | 2240.0 | 0.013393 | 0.114976 | 0.0 | 0.00 |
| Response | 2240.0 | 0.149107 | 0.356274 | 0.0 | 0.00 |
| Complain | 2240.0 | 0.009375 | 0.096391 | 0.0 | 0.00 |
| | | | | | |
| | 50% | 75% | max | | |
| ID | 5458.5 | 8427.75 | 11191.0 | | |
| Year_Birth | 1970.0 | 1977.00 | 1996.0 | | |
| Kidhome | 0.0 | 1.00 | 2.0 | | |

| | | | |
|---------------------|-------|--------|--------|
| Teenhome | 0.0 | 1.00 | 2.0 |
| Recency | 49.0 | 74.00 | 99.0 |
| MntWines | 173.5 | 504.25 | 1493.0 |
| MntFruits | 8.0 | 33.00 | 199.0 |
| MntMeatProducts | 67.0 | 232.00 | 1725.0 |
| MntFishProducts | 12.0 | 50.00 | 259.0 |
| MntSweetProducts | 8.0 | 33.00 | 263.0 |
| MntGoldProds | 24.0 | 56.00 | 362.0 |
| NumDealsPurchases | 2.0 | 3.00 | 15.0 |
| NumWebPurchases | 4.0 | 6.00 | 27.0 |
| NumCatalogPurchases | 2.0 | 4.00 | 28.0 |
| NumStorePurchases | 5.0 | 8.00 | 13.0 |
| NumWebVisitsMonth | 6.0 | 7.00 | 20.0 |
| AcceptedCmp3 | 0.0 | 0.00 | 1.0 |
| AcceptedCmp4 | 0.0 | 0.00 | 1.0 |
| AcceptedCmp5 | 0.0 | 0.00 | 1.0 |
| AcceptedCmp1 | 0.0 | 0.00 | 1.0 |
| AcceptedCmp2 | 0.0 | 0.00 | 1.0 |
| Response | 0.0 | 0.00 | 1.0 |
| Complain | 0.0 | 0.00 | 1.0 |

df.dtypes

| | |
|---------------------|--------|
| ID | int64 |
| Year_Birth | int64 |
| Education | object |
| Marital_Status | object |
| Income | object |
| Kidhome | int64 |
| Teenhome | int64 |
| Dt_Customer | object |
| Recency | int64 |
| MntWines | int64 |
| MntFruits | int64 |
| MntMeatProducts | int64 |
| MntFishProducts | int64 |
| MntSweetProducts | int64 |
| MntGoldProds | int64 |
| NumDealsPurchases | int64 |
| NumWebPurchases | int64 |
| NumCatalogPurchases | int64 |
| NumStorePurchases | int64 |
| NumWebVisitsMonth | int64 |
| AcceptedCmp3 | int64 |
| AcceptedCmp4 | int64 |
| AcceptedCmp5 | int64 |
| AcceptedCmp1 | int64 |
| AcceptedCmp2 | int64 |
| Response | int64 |
| Complain | int64 |

```
Country          object
dtype: object
```

- From the above description, we can confirm that there are some outliers in our Series

[Back to Table of Contents](#)

```
# Clean the column name ' Income ' by removing leading and trailing spaces
df.rename(columns=lambda x: x.strip(), inplace=True)
```

```
# Count missing values in the 'Income' column
missing_values_count = df['Income'].isnull().sum()
print("Total missing values in Income column:", missing_values_count)
```

```
Total missing values in Income column: 24
```

```
# Drop rows with missing values in the cleaned column
df.dropna(subset = 'Income', inplace=True)
assert(df['Income'].isnull().sum()==0,"Successfully removed missing values", "Failed to remove missing values")
```

```
# Assert that there are no missing values in the 'Income' column
assert df['Income'].isnull().sum() == 0, "Failed to remove missing values in 'Income' column"
print("Successfully removed missing values from 'Income' column")
```

```
Successfully removed missing values from 'Income' column
```

```
df.isna().sum()
```


| | |
|-------------------|---|
| ID | 0 |
| Year_Birth | 0 |
| Education | 0 |
| Marital_Status | 0 |
| Income | 0 |
| Kidhome | 0 |
| Teenhome | 0 |
| Dt_Customer | 0 |
| Recency | 0 |
| MntWines | 0 |
| MntFruits | 0 |
| MntMeatProducts | 0 |
| MntFishProducts | 0 |
| MntSweetProducts | 0 |
| MntGoldProds | 0 |
| NumDealsPurchases | 0 |


```

NumWebPurchases      0
NumCatalogPurchases  0
NumStorePurchases    0
NumWebVisitsMonth     0
AcceptedCmp3          0
AcceptedCmp4          0
AcceptedCmp5          0
AcceptedCmp1          0
AcceptedCmp2          0
Response              0
Complain              0
Country               0
dtype: int64

```

```
msno.bar(df);
```



- Here y axis ranges from 0 to 1, 1 means there is no missing values and 0 means all the data are missing. -So, if data ranges between 0 to 1 tells that the missing values are present in the dataset
- In our case, We can see that there is no Missing data anymore.

```

# Checking the number of duplicate values
print("Number of Duplicate values :", df.duplicated().sum())

```

```
Number of Duplicate values : 0
```

```
df.describe(include='object').T
```

| | count | unique | top | freq |
|----------------|-------|--------|------------|------|
| Education | 2216 | 5 | Graduation | 1116 |
| Marital_Status | 2216 | 8 | Married | 857 |

| | | | | |
|-------------|------|------|------------|------|
| Income | 2216 | 1974 | \$7,500.00 | 12 |
| Dt_Customer | 2216 | 662 | 8/31/12 | 12 |
| Country | 2216 | 8 | SP | 1093 |

- Here We can see, Income contains the numeric data, but because of dollar(\$) sign, it is treated as object, so lets preprocess it

```
df['income_in_usd'] =
df['Income'].str.replace(',','').str.replace('$','').astype(float)
```

- Now, we have customer's Income in proper format as income_in_usd. So, the previous Income feature is redundant and not useful anymore. This is good idea to drop it

```
#dropping redundant features
df.drop(columns='Income',inplace=True)
```

Checking % cover by each categorical labels

This is important because Rare values or labels in categorical variables can lead to overfitting in machine learning models, as they might not have sufficient representation in either the training or test set, causing instability in predictions

```
categorical = [var for var in df.columns if df[var].dtype=='O']
```

```
# check the number of different labels
```

```
for var in categorical:
    print(f"Columns Name : {var}")
    print()
    print(df[var].value_counts() / np.float(len(df)))
    print()
```

```
Columns Name : Education
```

| | |
|------------|----------|
| Graduation | 0.503610 |
| PhD | 0.217058 |
| Master | 0.164711 |
| 2n Cycle | 0.090253 |
| Basic | 0.024368 |

Name: Education, dtype: float64

```
Columns Name : Marital_Status
```

| | |
|----------|----------|
| Married | 0.386733 |
| Together | 0.258574 |
| Single | 0.212545 |
| Divorced | 0.104693 |
| Widow | 0.034296 |
| Alone | 0.001354 |
| YOL0 | 0.000903 |

```

Absurd      0.000903
Name: Marital_Status, dtype: float64

Columns Name : Dt_Customer

8/31/12      0.005415
2/14/13      0.004964
5/12/14      0.004964
9/12/12      0.004964
8/20/13      0.004513
...
1/9/14       0.000451
9/5/12       0.000451
11/9/13      0.000451
7/20/13      0.000451
9/1/12       0.000451
Name: Dt_Customer, Length: 662, dtype: float64

Columns Name : Country

SP      0.493231
SA      0.152076
CA      0.120036
AUS     0.066336
IND     0.066336
GER     0.052347
US      0.048285
ME      0.001354
Name: Country, dtype: float64

```

[Back to Table of Contents](#)

```

df.columns
Index(['ID', 'Year_Birth', 'Education', 'Marital_Status', 'Kidhome',
      'Teenhome', 'Dt_Customer', 'Recency', 'MntWines', 'MntFruits',
      'MntMeatProducts', 'MntFishProducts', 'MntSweetProducts',
      'MntGoldProds', 'NumDealsPurchases', 'NumWebPurchases',
      'NumCatalogPurchases', 'NumStorePurchases',
      'NumWebVisitsMonth',
      'AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5', 'AcceptedCmp1',
      'AcceptedCmp2', 'Response', 'Complain', 'Country',
      'income_in_usd'],
      dtype='object')

```

```
#check the datatypes of date column
df['Dt_Customer'].dtypes

dtype('O')
```

There can be Parsing issue as 'Dt_Customer' column which should be recognized as DateTime is now in Object format. So let's typecast it into correct format

```
df['Dt_Customer'] = pd.to_datetime(df['Dt_Customer'])

print("The newest customer's enrolment date in the records:",
max(df['Dt_Customer']))
print("The oldest customer's enrolment date in the records:",
min(df['Dt_Customer']))

The newest customer's enrolment date in the records: 2014-06-29
00:00:00
The oldest customer's enrolment date in the records: 2012-07-30
00:00:00
```

Since we have date of birth of each customer, we can get the respective age by reducing their dob by current year. For example, if customer A is born on 2000 and the current running year is 2020, the age of that individual is $2020 - 2000 = 20$ years. As we have the record from 2012 to 2014, we will consider the age of customer by 2015.

```
# As we have the record from 2012 to 2014, we will consider the age of
customer by 2015.
df['Age'] = 2015 - df['Year_Birth']
```

Let's assume that we want to focus on Veg/Vegan Product only. So let's drop unnecessary features

```
df.drop(columns=['MntMeatProducts', 'MntFishProducts'], inplace=True)
```

Creating an additional attribute named "total_spent" that represents the overall expenditure made by the customer across different categories during the course of two years.

```
df['total_spent'] = df['MntWines'] + df['MntFruits'] +
df['MntSweetProducts'] + df['MntGoldProds']
```

Create another feature "Living_With" out of "Marital_Status" to extract the living situation of Customer.

```
df['Marital_Status'].value_counts()

Married      857
Together     573
Single       471
```

```

Divorced      232
Widow         76
Alone         3
YOLO          2
Absurd        2
Name: Marital_Status, dtype: int64

df['Living_With'] = df['Marital_Status'].replace({'Married':'Partner',
'Together':'Partner', 'Absurd':'Alone', 'Widow':'Alone',
'YOLO':'Alone', 'Divorced':'Alone', 'Single':'Alone'})

```

Recall That :

- Kidhome: Number of children in customer's household
- Teenhome: Number of teenagers in customer's household

So, Lets Create a feature "Children" to indicate total children in a household

```
df['Children'] = df['Kidhome'] + df['Teenhome']
```

To gain a better insight into the structure of the household, it's recommended to create a new feature that reflects the "Family_Size."

```
df['Family_Size'] = df['Living_With'].replace({'Alone': 1,
'Partner':2}) + df['Children']
```

An additional idea is to incorporate a feature named "Is_Parent," which would serve the purpose of indicating whether someone is a parent or not.

```
df['Is_Parent'] = np.where(df.Children > 0, 1, 0)
```

Grouping education levels in three categories

```
df['Education'] = df['Education'].replace({'Basic':'Undergraduate',
'2n Cycle':'Undergraduate', 'Graduation':'Graduate',
'Master':'Postgraduate', 'PhD':'Postgraduate'})
```

Creating new column accepted_promotion to store the number of promotion campaign accepted by customer

```
df['accepted_promotion'] = df['AcceptedCmp1'] + df['AcceptedCmp2'] +
df['AcceptedCmp3'] + df['AcceptedCmp4'] + df['AcceptedCmp5']

df.head()
```

| | ID | Year_Birth | Education | Marital_Status | Kidhome | Teenhome |
|---------------|------------|------------|-----------|----------------|---------|----------|
| Dt_Customer \ | | | | | | |
| 0 | 1826 | 1970 | Graduate | Divorced | 0 | 0 |
| | 2014-06-16 | | | | | |

| | | | | | | |
|------------|-------|------|----------|----------|---|---|
| 1 | 1 | 1961 | Graduate | Single | 0 | 0 |
| 2014-06-15 | | | | | | |
| 2 | 10476 | 1958 | Graduate | Married | 0 | 1 |
| 2014-05-13 | | | | | | |
| 3 | 1386 | 1967 | Graduate | Together | 1 | 1 |
| 2014-05-11 | | | | | | |
| 4 | 5371 | 1989 | Graduate | Single | 1 | 0 |
| 2014-04-08 | | | | | | |

| | Recency | MntWines | MntFruits | MntSweetProducts | MntGoldProds | \ |
|---|---------|----------|-----------|------------------|--------------|---|
| 0 | 0 | 189 | 104 | 189 | 218 | |
| 1 | 0 | 464 | 5 | 0 | 37 | |
| 2 | 0 | 134 | 11 | 2 | 30 | |
| 3 | 0 | 10 | 0 | 0 | 0 | |
| 4 | 0 | 6 | 16 | 0 | 34 | |

| | NumDealsPurchases | NumWebPurchases | NumCatalogPurchases |
|---------------------|-------------------|-----------------|---------------------|
| NumStorePurchases \ | | | |
| 0 | 1 | 4 | 4 |
| 6 | | | |
| 1 | 1 | 7 | 3 |
| 7 | | | |
| 2 | 1 | 3 | 2 |
| 5 | | | |
| 3 | 1 | 1 | 0 |
| 2 | | | |
| 4 | 2 | 3 | 1 |
| 2 | | | |

| | NumWebVisitsMonth | AcceptedCmp3 | AcceptedCmp4 | AcceptedCmp5 |
|----------------|-------------------|--------------|--------------|--------------|
| AcceptedCmp1 \ | | | | |
| 0 | 1 | 0 | 0 | 0 |
| 0 | | | | |
| 1 | 5 | 0 | 0 | 0 |
| 0 | | | | |
| 2 | 2 | 0 | 0 | 0 |
| 0 | | | | |
| 3 | 7 | 0 | 0 | 0 |
| 0 | | | | |
| 4 | 7 | 1 | 0 | 0 |
| 0 | | | | |

| | AcceptedCmp2 | Response | Complain | Country | income_in_usd | Age |
|---------------|--------------|----------|----------|---------|---------------|-----|
| total_spent \ | | | | | | |
| 0 | 0 | 1 | 0 | SP | 84835.0 | 45 |
| 700 | | | | | | |
| 1 | 1 | 1 | 0 | CA | 57091.0 | 54 |
| 506 | | | | | | |
| 2 | 0 | 0 | 0 | US | 67267.0 | 57 |
| 177 | | | | | | |

| 3 | 0 | 0 | 0 | AUS | 32474.0 | 48 |
|----|-------------|----------|-------------|-----------|--------------------|----|
| 10 | | | | | | |
| 4 | 0 | 1 | 0 | SP | 21474.0 | 26 |
| 56 | | | | | | |
| | Living_With | Children | Family_Size | Is_Parent | accepted_promotion | |
| 0 | Alone | 0 | 1 | 0 | | 0 |
| 1 | Alone | 0 | 1 | 0 | | 1 |
| 2 | Partner | 1 | 3 | 1 | | 0 |
| 3 | Partner | 2 | 4 | 1 | | 0 |
| 4 | Alone | 1 | 2 | 1 | | 1 |

Dropping Redundant Features

- We made a new column using information we already had. This means some of the old columns are no longer needed, so we can get rid of them now. However, we will still keep some column as it can be useful for segmentation

```
col = (['Marital_Status', 'Dt_Customer', 'Year_Birth',
        'ID', 'AcceptedCmp1', 'AcceptedCmp2', 'AcceptedCmp3',
        'AcceptedCmp4', 'AcceptedCmp5', 'Kidhome', 'Teenhome'])
df.drop(columns=col, inplace=True)
df.tail()
```

| | Education | Recency | MntWines | MntFruits | MntSweetProducts | \ |
|------|---------------|---------|----------|-----------|------------------|-----|
| 2235 | Postgraduate | 99 | 372 | 18 | | 48 |
| 2236 | Undergraduate | 99 | 5 | 10 | | 8 |
| 2237 | Graduate | 99 | 185 | 2 | | 5 |
| 2238 | Graduate | 99 | 267 | 38 | | 165 |
| 2239 | Postgraduate | 99 | 169 | 24 | | 0 |


| | MntGoldProds | NumDealsPurchases | NumWebPurchases |
|---------------------|--------------|-------------------|-----------------|
| NumCatalogPurchases | \ | | |
| 2235 | 78 | 2 | 5 |
| 2 | | | |
| 2236 | 16 | 1 | 1 |
| 0 | | | |
| 2237 | 14 | 2 | 6 |
| 1 | | | |
| 2238 | 63 | 1 | 5 |
| 4 | | | |
| 2239 | 144 | 1 | 8 |
| 5 | | | |

| | NumStorePurchases | NumWebVisitsMonth | Response | Complain | Country |
|------|-------------------|-------------------|----------|----------|---------|
| \ | | | | | |
| 2235 | 11 | 4 | 0 | 0 | US |
| 2236 | 3 | 8 | 0 | 0 | SP |

| | | | | | |
|---------------|---------------|--------------------|-------------|-------------|----------|
| 2237 | 5 | 8 | 0 | 0 | SP |
| 2238 | 10 | 3 | 0 | 0 | IND |
| 2239 | 4 | 7 | 1 | 0 | CA |
| | | | | | |
| | income_in_usd | Age | total_spent | Living_With | Children |
| Family_Size \ | | | | | |
| 2235 | 66476.0 | 39 | 516 | Alone | 1 |
| 2 | | | | | |
| 2236 | 31056.0 | 38 | 39 | Partner | 1 |
| 3 | | | | | |
| 2237 | 46310.0 | 39 | 206 | Alone | 1 |
| 2 | | | | | |
| 2238 | 65819.0 | 37 | 533 | Partner | 0 |
| 2 | | | | | |
| 2239 | 94871.0 | 46 | 337 | Partner | 2 |
| 4 | | | | | |
| | | | | | |
| | Is_Parent | accepted_promotion | | | |
| 2235 | 1 | 0 | | | |
| 2236 | 1 | 0 | | | |
| 2237 | 1 | 0 | | | |
| 2238 | 0 | 0 | | | |
| 2239 | 1 | 2 | | | |

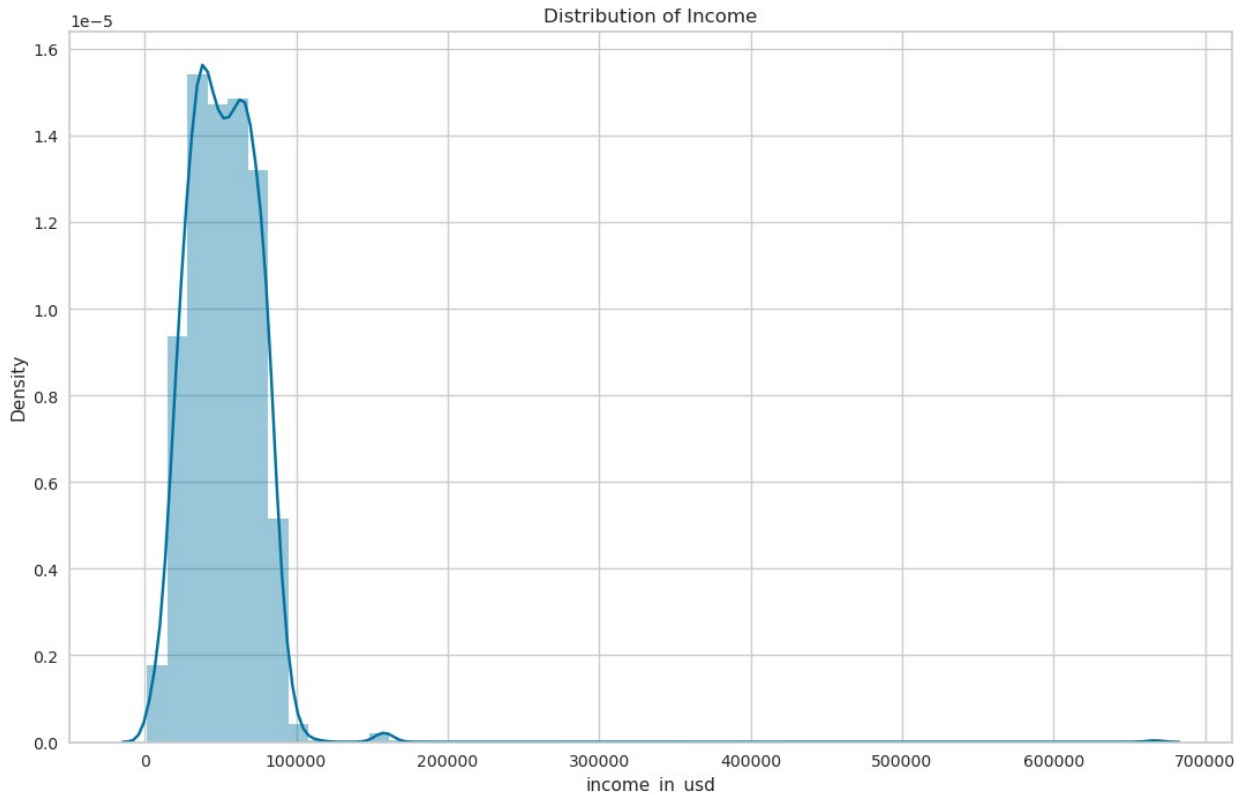
We can see that, the four feature has been successfully removed

[Back to Table of Contents](#)



Column : income_in_usd

```
#lets visualize the distribution  
plt.figure(figsize=(13,8))  
sns.distplot(df.income_in_usd);  
plt.title("Distribution of Income");
```



We can see a small peak going from 150,000 to 700,000. This shows the sign of rare/outlier values. And as the distribution of data is skewed towards the right, we will be using trimming techniques instead of capping for handling outliers.

```
higher_quantile_income = df['income_in_usd'].quantile(0.95)
lower_quantile_income = df['income_in_usd'].quantile(0.05)
print(f"Higher Quantile is {higher_quantile_income}")
print(f"Lower Quantile is {lower_quantile_income}")
print(f"There are {(df['income_in_usd'] >
higher_quantile_income).sum()} data above higher quantile")
print(f"There are {(df['income_in_usd'] <
lower_quantile_income).sum()} data less than lower quantile")
```

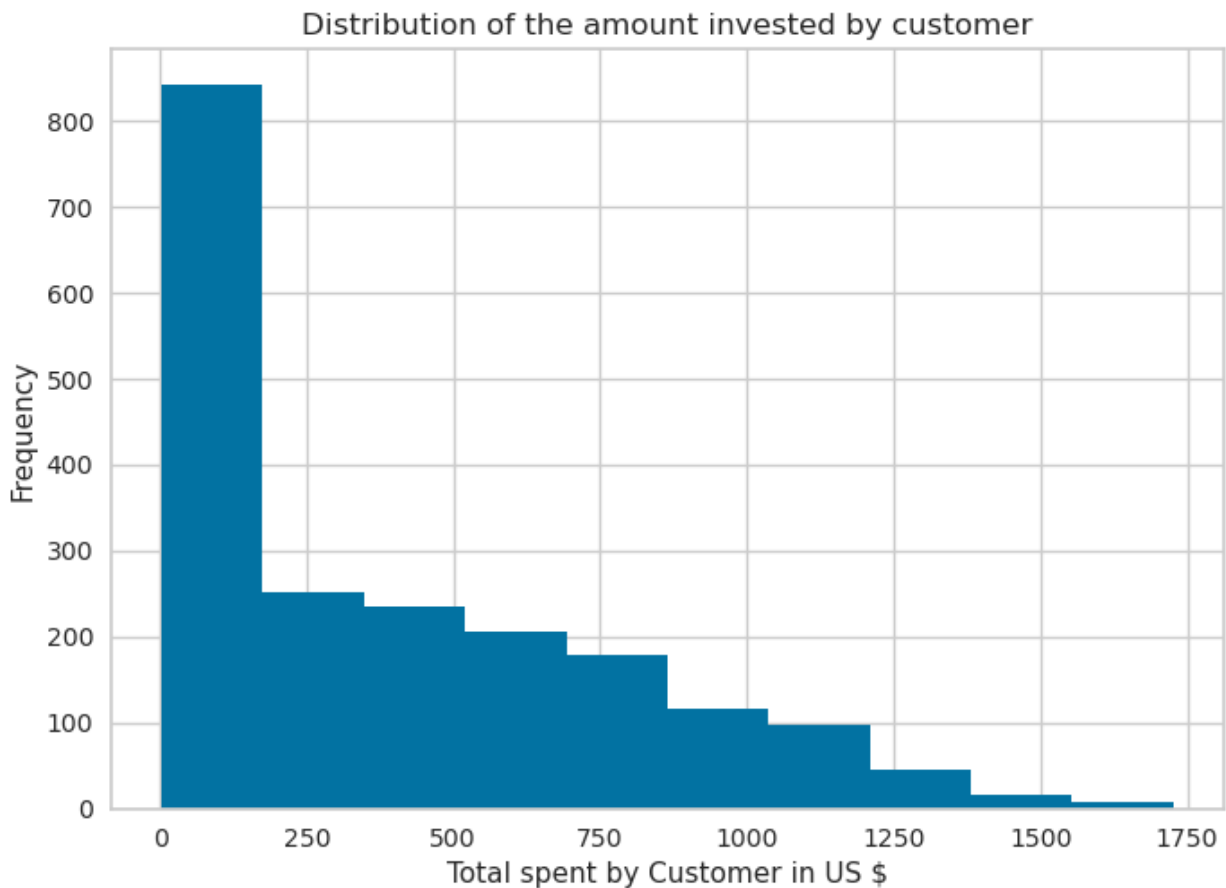
```
Higher Quantile is 84130.0
Lower Quantile is 18985.5
There are 111 data above higher quantile
There are 111 data less than lower quantile
```

```
#creating a mask of trimmed income
mask = (df['income_in_usd'] <= higher_quantile_income) &
(df['income_in_usd'] >= lower_quantile_income)
df = df[mask]
print(f"Maximum value after capping: {df['income_in_usd'].max()}")
print(f"Minimum value after capping: {df['income_in_usd'].min()}")
```

Maximum value after capping: 84117.0
Minimum value after capping: 18988.0

Column : total_spent

```
#plotting histogram of total_spent of customer  
plt.hist(df['total_spent'])  
plt.xlabel("Total spent by Customer in US $")  
plt.ylabel("Frequency");  
plt.title("Distribution of the amount invested by customer");
```



As the distribution is highly left skewed, lets analyze it and perform trimming

```
higher_quantile_total_spent = df['total_spent'].quantile(0.95)  
lower_quantile_total_spent = df['total_spent'].quantile(0.05)  
print(f"Higher Quantile is {higher_quantile_total_spent}")  
print(f"Lower Quantile is {lower_quantile_total_spent}")  
print()  
print(f"There are {(df['total_spent'] >  
higher_quantile_total_spent).sum()} data above higher quantile")  
print(f"There are {(df['total_spent'] <  
lower_quantile_total_spent).sum()} data less than lower quantile")
```

```
#lets trim/remove th outside quantile data
updated_df = df[(df['total_spent'] >= lower_quantile_total_spent) &
(df['total_spent'] <= higher_quantile_total_spent)]
```

```
print(f"Maximum spending value after trimming:
{updated_df['total_spent'].max()}")
print(f"Minimum spending value after trimming:
{updated_df['total_spent'].min()}")
```

Higher Quantile is 1145.0
Lower Quantile is 15.0

There are 99 data above higher quantile
There are 94 data less than lower quantile
Maximum spending value after trimming: 1145
Minimum spending value after trimming: 15

Comparing Shape of updated dataframe with previous


```
print("Shape before updated {}, Shape after updated
{}".format(df.shape, updated_df.shape))
```

Shape before updated (1994, 22), Shape after updated (1801, 22)


Column : Age

```
updated_df["Age"].min(), updated_df['Age'].max()
(20, 115)
```

We can see that minimum age is Okey, but Maximum age is too far that normal customer can have. So, in this senario IQR methods seems better option which we see in next steps



```
#lets visualize the distribution
plt.figure(figsize=(13,8))
sns.distplot(df.Age,color='purple');
plt.title("Distribution of Age");
```



```

# Calculate the interquartile range (IQR)
Q1 = updated_df['Age'].quantile(0.25)
Q3 = updated_df['Age'].quantile(0.75)
IQR = Q3 - Q1

# Define lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
print(f"lower_bound :{lower_bound}, upper_bound :{upper_bound}")
print()

# Identify and print the number of outliers
outliers_below = (updated_df['Age'] < lower_bound).sum()
outliers_above = (updated_df['Age'] > upper_bound).sum()
print(f"Number of outliers below the lower bound: {outliers_below}")
print(f"Number of outliers above the upper bound: {outliers_above}")

updated_df[(updated_df['Age']>=upper_bound) |
(updated_df['Age']<=lower_bound)]

lower_bound :13.5, upper_bound :81.5

Number of outliers below the lower bound: 0
Number of outliers above the upper bound: 1

```

| | | | | | | |
|------|---------------|---------|----------|-----------|------------------|---|
| | Education | Recency | MntWines | MntFruits | MntSweetProducts | \ |
| 2233 | Undergraduate | 99 | 15 | 6 | | 4 |

| | | | |
|---------------------|--------------|-------------------|-----------------|
| | MntGoldProds | NumDealsPurchases | NumWebPurchases |
| NumCatalogPurchases | \ | | |
| 2233 | 25 | 1 | 2 |

| | | | | | |
|------|-------------------|-------------------|----------|----------|---------|
| | NumStorePurchases | NumWebVisitsMonth | Response | Complain | Country |
| \ | | | | | |
| 2233 | 2 | 5 | 0 | 1 | IND |

| | | | | | |
|-------------|---------------|-----|-------------|-------------|----------|
| | income_in_usd | Age | total_spent | Living_With | Children |
| Family_Size | \ | | | | |
| 2233 | 36640.0 | 115 | 50 | Alone | 1 |

| | | |
|------|-----------|--------------------|
| | Is_Parent | accepted_promotion |
| 2233 | 1 | 0 |

We can see there is only one data to be removed

```

# # Remove outliers using IQR technique
updated_df = df[(df['Age'] >= lower_bound) & (df['Age'] <=

```

```
upper_bound)]
```

```
# Print the maximum and minimum age values after outlier removal  
print(f"Maximum age value after outlier removal:  
{updated_df['Age'].max()}")  
print(f"Minimum age value after outlier removal:  
{updated_df['Age'].min()}")
```

Maximum age value after outlier removal: 75

Minimum age value after outlier removal: 20

Visualizing through boxplot after handling outliers


```
# Create boxplots for all three columns  
plt.figure(figsize=(18, 6))
```

```
# Boxplot for 'income_in_usd'  
plt.subplot(1, 3, 1)  
plt.boxplot(updated_df['income_in_usd'])  
plt.xlabel("Income in USD")  
plt.ylabel("Values")  
plt.title("Boxplot of Income After Handling Outliers")
```

```
# Boxplot for 'total_spent'  
plt.subplot(1, 3, 2)  
plt.boxplot(updated_df['total_spent'])  
plt.xlabel("Total Spent by Customer")  
plt.ylabel("Values")  
plt.title("Boxplot of Total Spent After Handling Outliers")
```

```
# Boxplot for 'Age'  
plt.subplot(1, 3, 3)  
plt.boxplot(updated_df['Age'])  
plt.xlabel("Age")  
plt.ylabel("Values")  
plt.title("Boxplot of Age After Handling Outliers")
```

```
plt.tight_layout()  
plt.show()
```



[Back to Table of Contents](#)


Let's check the most spending customer by Country

```
# plotting the pie chart
fig = px.pie(
    updated_df,
    values = 'total_spent',
    names = 'Country',
    title = 'Total Spent by Country',
    color_discrete_sequence = px.colors.sequential.Magma
)

fig.show()
```


After analyzing the pie chart depicting the distribution of total spending by country label, several key insights emerge. Spain emerges as the dominant contributor, accounting for a significant portion of the spending at 49%. Following behind, South Africa constitutes 15.7% of the total expenditure, while Australia contributes 12.7%. India and Germany also play noteworthy roles, with contributions of 6.36% and 5.78%, respectively. The United States and Mexico, although smaller in comparison, still hold significance, representing 5.02% and 0.232% of the total spend, respectively. These insights shed light on the distribution of expenditures across various countries, allowing for informed decision-making and strategic planning.

```
plt.figure(figsize=(13,8))
sns.scatterplot(x=updated_df['income_in_usd'],
y=updated_df['total_spent'], color='#cc0000')
plt.title("Income vs total spent of a Customer");
```

We can see a positive linear relationship of income and expenditure of a customer

```
plt.figure(figsize=(13,8))
sns.scatterplot(x=updated_df['Age'], y=updated_df['total_spent'])
plt.title("Age vs Total spent of customer");
```



```
fig = px.scatter(updated_df,
                 x='Education',
                 y='total_spent',
                 color='Education',
                 title='Relationship between Education and Total
Spent',
                 labels={'education': 'Education', 'total_spent':
'Total Spent'})

# Show the plot
fig.show()
```

We can conclude, more the person is educationally qualified, more he/she spends

```
# Count the occurrences of each education level
education_counts = updated_df['Education'].value_counts()

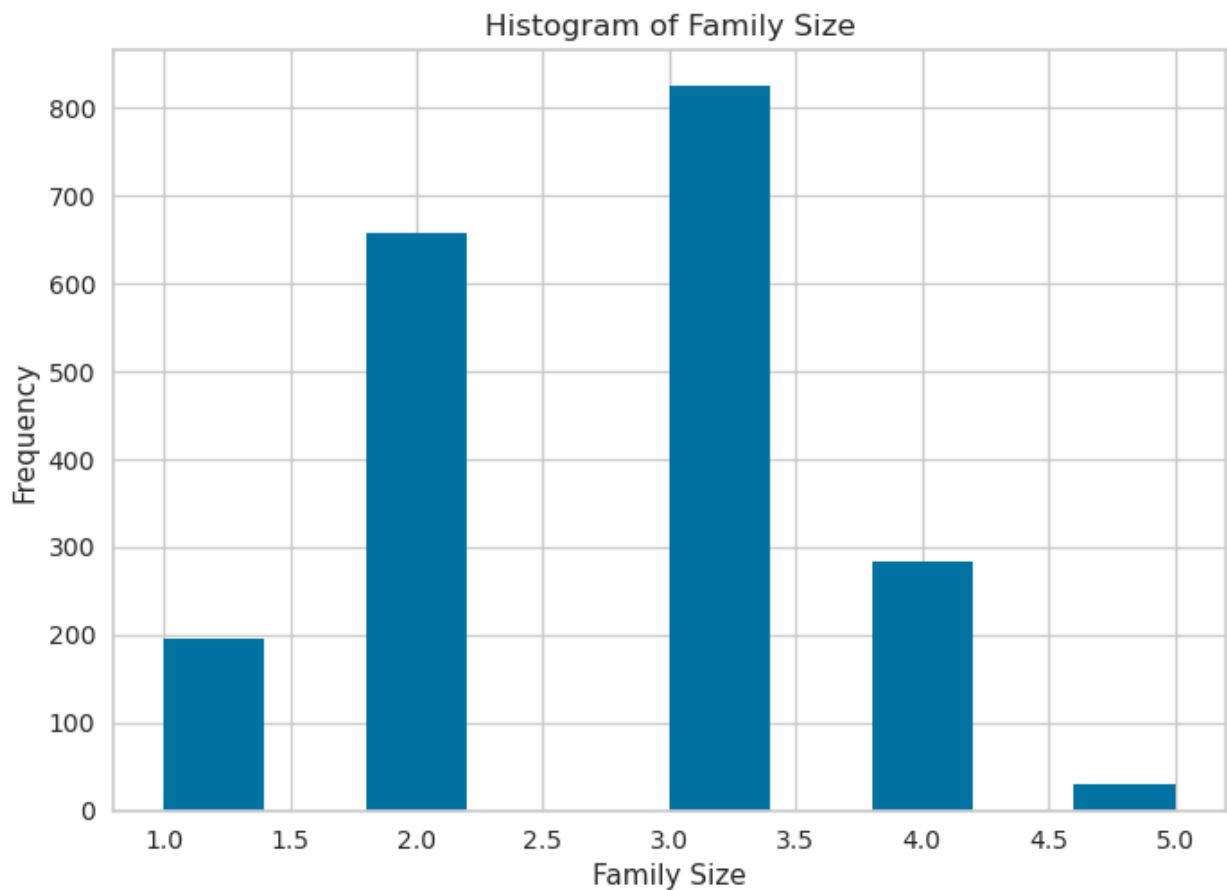
# Create a pie chart using Plotly
fig = px.pie(
    education_counts,
    names=education_counts.index,
    values=education_counts.values,
    title='Distribution of Education Levels',
)
```

```

# Show the pie chart
fig.show()

updated_df['Family_Size'].plot(
    kind='hist',
    title='Histogram of Family Size'
)
plt.xlabel('Family Size')
plt.ylabel('Frequency')
plt.show()

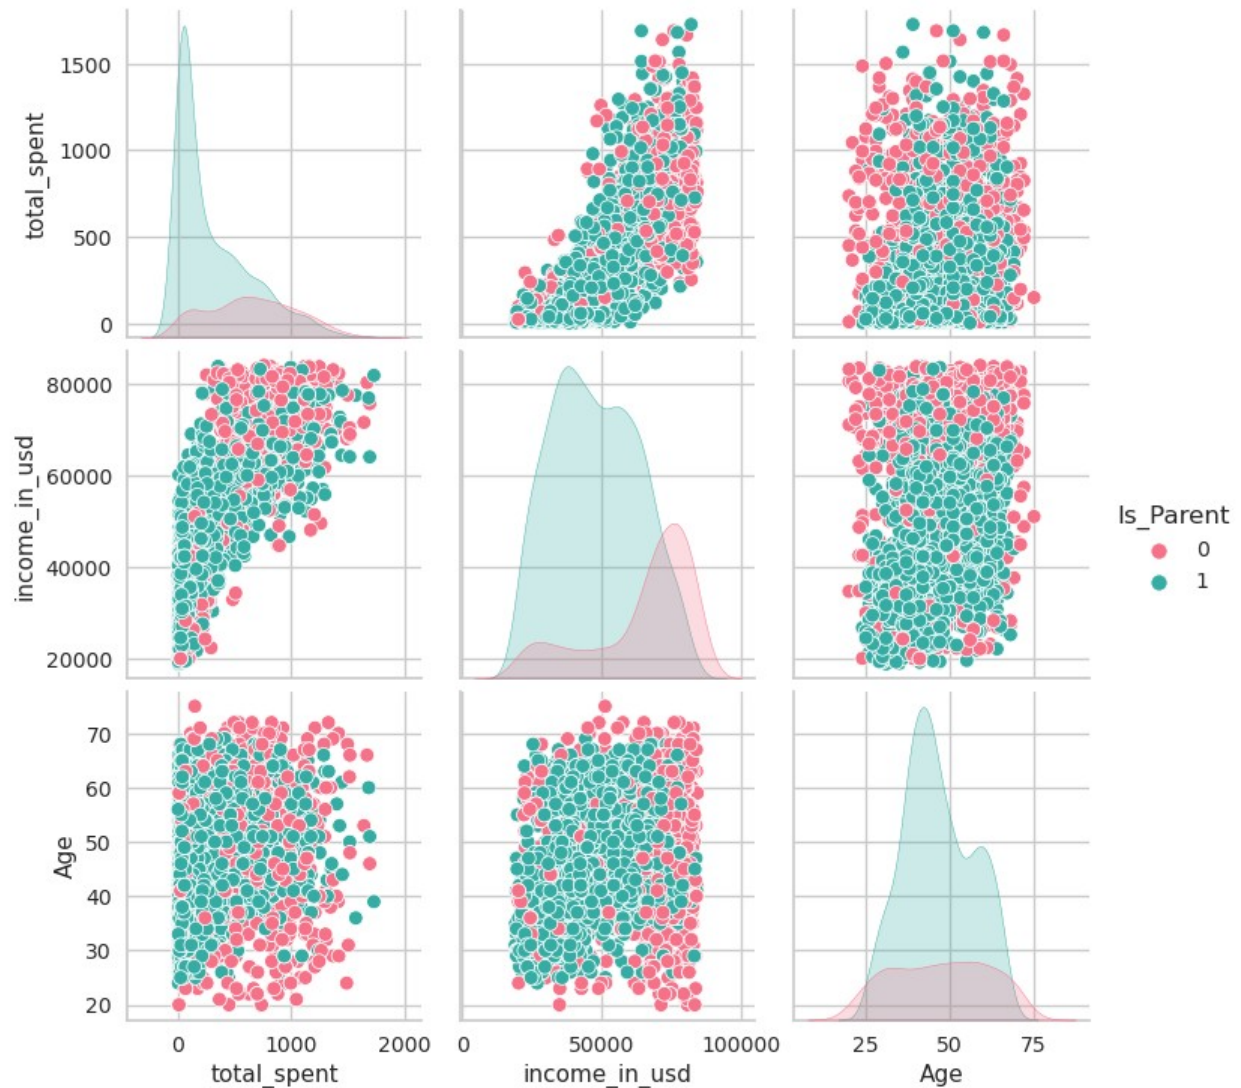
```



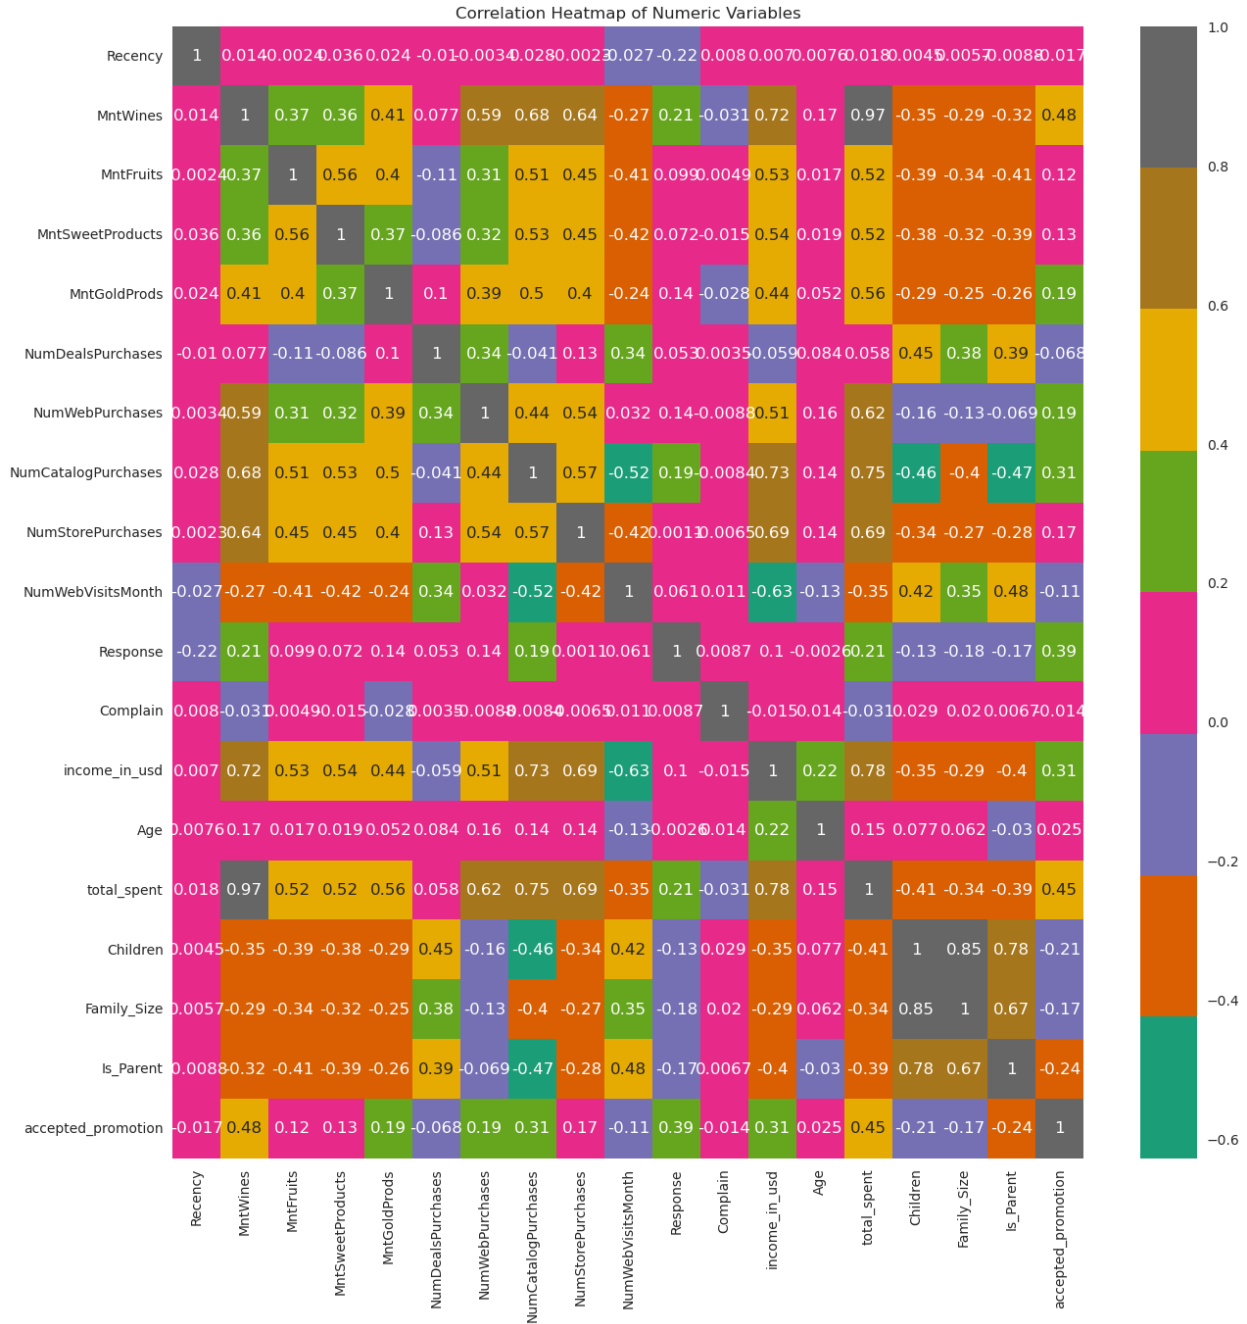
```

#plotting the pairplot
sns.pairplot(
    updated_df ,
    vars=['total_spent', 'income_in_usd', 'Age'] ,
    hue='Is_Parent',
    palette='husl'
);

```



```
# Create the correlation heatmap using Seaborn
plt.figure(figsize=(15,15))
sns.heatmap(updated_df.corr(), annot=True, cmap='Dark2')
plt.title('Correlation Heatmap of Numeric Variables')
plt.show()
```

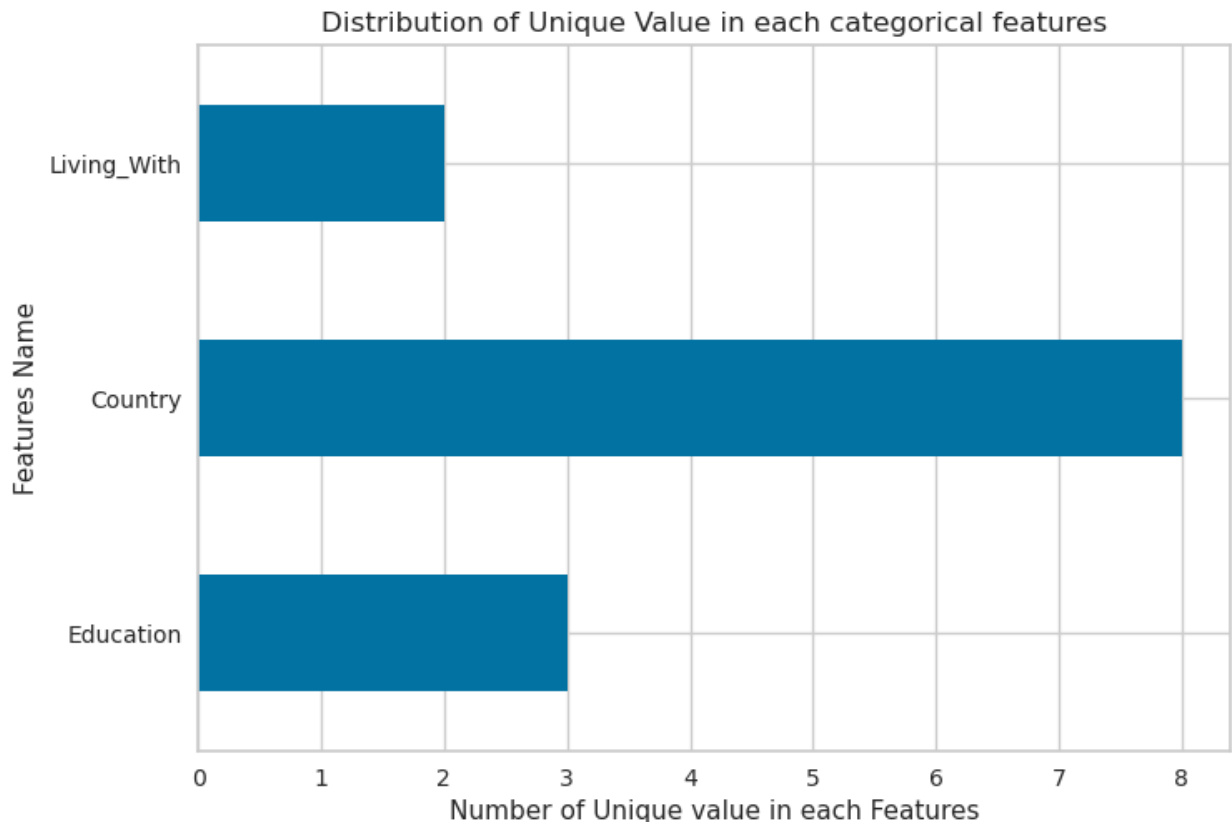


It seems like we have correlation among the feature. It increases the data redundancy and effect in further evaluation. So it is best for reducing the dimensionality

[Back to Table of Contents](#)

```
categorical_data = [col for col in updated_df.columns if
updated_df[col].dtype=='O']
categorical_data
```

```
['Education', 'Country', 'Living_With']
updated_df[categorical_data].nunique().plot(
    kind='barh',
    xlabel="Number of Unique value in each Features",
    ylabel="Features Name",
    title="Distribution of Unique Value in each categorical
features"
);
```



From above horizontal bar plot, we can see that there are 2 Unique value in **Living_With** column, 8 unique value in **Country** column and 3 unique value in **Education** column

```
# Map categorical values to numerical values in 'Education' column
updated_df['Education'] =
updated_df['Education'].map({'Undergraduate': 0, 'Graduate': 1,
'Postgraduate': 2})

# Map categorical values to numerical values in 'Living_With' column
updated_df['Living_With'] = updated_df['Living_With'].map({'Alone': 0,
'Partner': 1})

updated_df['Country'].value_counts().to_dict()
```

```
{'SP': 975,
'SA': 305,
'CA': 237,
'AUS': 136,
'IND': 130,
'GER': 109,
'US': 96,
'ME': 3}
```

Frequency Encoding - Encode the countries based on their frequency, where higher frequency countries get higher encoded values. Countries with higher frequencies will have higher encoded values, and those with lower frequencies will have lower encoded values. This approach can help capture the relationship between countries and the total spent more accurately.

```
country_frequency = updated_df['Country'].value_counts().to_dict()
updated_df['Country'] = updated_df['Country'].map(country_frequency)
updated_df.head()
```

| | Education | Recency | MntWines | MntFruits | MntSweetProducts |
|----------------|-----------|---------|----------|-----------|------------------|
| MntGoldProds \ | | | | | |
| 1 | 1 | 0 | 464 | 5 | 0 |
| 37 | | | | | |
| 2 | 1 | 0 | 134 | 11 | 2 |
| 30 | | | | | |
| 3 | 1 | 0 | 10 | 0 | 0 |
| 0 | | | | | |
| 4 | 1 | 0 | 6 | 16 | 0 |
| 34 | | | | | |
| 5 | 2 | 0 | 336 | 130 | 32 |
| 43 | | | | | |

| | NumDealsPurchases | NumWebPurchases | NumCatalogPurchases |
|---------------------|-------------------|-----------------|---------------------|
| NumStorePurchases \ | | | |
| 1 | 1 | 7 | 3 |
| 7 | | | |
| 2 | 1 | 3 | 2 |
| 5 | | | |
| 3 | 1 | 1 | 0 |
| 2 | | | |
| 4 | 2 | 3 | 1 |
| 2 | | | |
| 5 | 1 | 4 | 7 |
| 5 | | | |

| | NumWebVisitsMonth | Response | Complain | Country | income_in_usd | Age |
|---|-------------------|----------|----------|---------|---------------|-----|
| \ | | | | | | |
| 1 | 5 | 1 | 0 | 237 | 57091.0 | 54 |

| | | | | | | |
|---|---|---|---|-----|---------|----|
| 2 | 2 | 0 | 0 | 96 | 67267.0 | 57 |
| 3 | 7 | 0 | 0 | 136 | 32474.0 | 48 |
| 4 | 7 | 1 | 0 | 975 | 21474.0 | 26 |
| 5 | 2 | 1 | 0 | 975 | 71691.0 | 57 |

| | total_spent | Living_With | Children | Family_Size | Is_Parent | \ |
|---|-------------|-------------|----------|-------------|-----------|---|
| 1 | 506 | 0 | 0 | 1 | 0 | |
| 2 | 177 | 1 | 1 | 3 | 1 | |
| 3 | 10 | 1 | 2 | 4 | 1 | |
| 4 | 56 | 0 | 1 | 2 | 1 | |
| 5 | 541 | 0 | 0 | 1 | 0 | |

| | accepted_promotion |
|---|--------------------|
| 1 | 1 |
| 2 | 0 |
| 3 | 0 |
| 4 | 1 |
| 5 | 0 |

[Back to Table of Contents](#)

```
old_df = updated_df.copy()
```

Standard Scaling

```
ss = StandardScaler()
data = pd.DataFrame(ss.fit_transform(updated_df), columns =
updated_df.columns)
```

```
data.head()
```

| | Education | Recency | MntWines | MntFruits | MntSweetProducts | |
|--------------|-----------|-----------|-----------|-----------|------------------|-----------|
| MntGoldProds | \ | | | | | |
| 0 | -0.448852 | -1.700523 | 0.507860 | -0.530234 | -0.657691 | -0.139532 |
| 1 | -0.448852 | -1.700523 | -0.514270 | -0.375133 | -0.606700 | -0.276239 |
| 2 | -0.448852 | -1.700523 | -0.898343 | -0.659485 | -0.657691 | -0.862127 |
| 3 | -0.448852 | -1.700523 | -0.910732 | -0.245882 | -0.657691 | -0.198121 |
| 4 | 1.113499 | -1.700523 | 0.111398 | 2.701040 | 0.158158 | -0.022355 |

| | NumDealsPurchases | NumWebPurchases | NumCatalogPurchases | | | |
|---------------------|--------------------|-----------------|---------------------|-------------|---------------|---|
| NumStorePurchases \ | | | | | | |
| 0 | -0.749126 | 1.084025 | 0.162801 | | | |
| 0.347899 | | | | | | |
| 1 | -0.749126 | -0.428444 | -0.213228 | - | | |
| 0.275283 | | | | | | |
| 2 | -0.749126 | -1.184679 | -0.965287 | - | | |
| 1.210056 | | | | | | |
| 3 | -0.205375 | -0.428444 | -0.589258 | - | | |
| 1.210056 | | | | | | |
| 4 | -0.749126 | -0.050327 | 1.666920 | - | | |
| 0.275283 | | | | | | |
| | NumWebVisitsMonth | Response | Complain | Country | income_in_usd | |
| Age \ | | | | | | |
| 0 | -0.162920 | 2.524693 | -0.095515 | -0.879956 | 0.299843 | |
| 0.645898 | | | | | | |
| 1 | -1.518239 | -0.396088 | -0.095515 | -1.240873 | 0.872273 | |
| 0.907934 | | | | | | |
| 2 | 0.740627 | -0.396088 | -0.095515 | -1.138485 | -1.084936 | |
| 0.121827 | | | | | | |
| 3 | 0.740627 | 2.524693 | -0.095515 | 1.009098 | -1.703718 | |
| 1.799767 | | | | | | |
| 4 | -1.518239 | 2.524693 | -0.095515 | 1.009098 | 1.121136 | |
| 0.907934 | | | | | | |
| | total_spent | Living_With | Children | Family_Size | Is_Parent | \ |
| 0 | 0.288532 | -1.365558 | -1.342302 | -1.838700 | -1.719957 | |
| 1 | -0.570440 | 0.732301 | 0.006776 | 0.395570 | 0.581410 | |
| 2 | -1.006452 | 0.732301 | 1.355854 | 1.512705 | 0.581410 | |
| 3 | -0.886353 | -1.365558 | 0.006776 | -0.721565 | 0.581410 | |
| 4 | 0.379911 | -1.365558 | -1.342302 | -1.838700 | -1.719957 | |
| | accepted_promotion | | | | | |
| 0 | 1.288226 | | | | | |
| 1 | -0.427398 | | | | | |
| 2 | -0.427398 | | | | | |
| 3 | 1.288226 | | | | | |
| 4 | -0.427398 | | | | | |

Dimensionality Reduction

Principal component analysis(PCA)


The principal components of a collection of points in a real coordinate space are a sequence of p unit vectors, where the i -th vector is the direction of a line that best fits the data while being orthogonal to the first $i-1$ vectors. Here, a best-fitting line is defined as one that minimizes the average squared distance from the points to the line. These directions constitute an orthonormal basis in which different individual dimensions of the data are linearly uncorrelated.


Principal component analysis (PCA) is the process of computing the principal components and using them to perform a change of basis on the data, sometimes using only the first few principal components and ignoring the rest.

```
# Fit PCA to your data
pca = PCA()
pca.fit(data)

# Plot explained variance ratio
plt.plot(range(1, len(data.columns) + 1),
pca.explained_variance_ratio_.cumsum(), marker='o')
plt.xlabel('Number of Principal Components')
plt.ylabel('Explained Variance Ratio')
plt.title('Explained Variance Ratio vs. Number of Principal
Components')
plt.show()

# Scree plot
plt.plot(range(1, len(data.columns) + 1), pca.explained_variance_,
marker='o')
plt.xlabel('Principal Component Index')
plt.ylabel('Eigenvalue')
plt.title('Scree Plot')
plt.show()
```





The above second plot indicates that the instant falling of plot at Principal component Index of 3 with Eigen value of 1.5 is better choice

```
pca = PCA(n_components=3) #also more than 3 cluster is difficult for visualization
```

```
# Fit and transform your data
```

```
pca_result = pca.fit_transform(data) # X is your data matrix
```

```
# The transformed data with 3 principal components
```

```
print(pca_result)
```

```
[[ 2.15835178 -1.52409028  3.47087126]
 [-0.54788839 -0.37486955 -1.03647949]
 [-3.49630643 -0.04324013 -0.25435839]
 ...
 [-2.8049115  -1.24349908 -0.90581971]
 [-1.27017962 -0.47380907  0.46169894]
 [ 3.06333439 -1.62261756 -2.32542662]]
```

```
# Get the indices of the original columns that contribute most to the top 3 principal components
```

```
top_pca_indices = pca.components_.argsort(axis=1)[:,-3:] # Indices
of top 3 components for each feature
```

```
# Get the corresponding column names for the top 3 principal
components
```

```
top_pca_column_names = data.columns[top_pca_indices]
```

```
print("Top Principal Component Column Names:", top_pca_column_names)
```

```
Top Principal Component Column Names: [['NumCatalogPurchases'
'income_in_usd' 'total_spent']
['Children' 'Family_Size' 'NumDealsPurchases']
['NumWebVisitsMonth' 'accepted_promotion' 'Response']]
```

```
# Create a new DataFrame with transformed data
```

```
pca_df = pd.DataFrame(data=pca_result, columns=['PC1_mntwines',
'PC2_income', 'PC3_spent'])
```

```
# Display the new DataFrame
```

```
pca_df.head()
```

| | PC1_mntwines | PC2_income | PC3_spent |
|---|--------------|------------|-----------|
| 0 | 2.158352 | -1.524090 | 3.470871 |
| 1 | -0.547888 | -0.374870 | -1.036479 |
| 2 | -3.496306 | -0.043240 | -0.254358 |
| 3 | -1.833339 | -1.462887 | 3.689209 |
| 4 | 3.511079 | -2.253159 | 1.735582 |

```
# Extract columns for plotting
```

```
x = pca_df['PC1_mntwines']
```

```
y = pca_df['PC2_income']
```

```
z = pca_df['PC3_spent']
```

```
# Create a 3D scatter plot
```

```
fig = plt.figure(figsize=(20,15))
```

```
ax = fig.add_subplot(111, projection='3d')
```

```
ax.scatter(x, y, z, c='darkred', marker='*', label='Data Points')
```

```
ax.set_title('A 3D Projection of Data In the Reduced Dimension')
```

```
ax.set_xlabel('PC1')
```


```
ax.set_ylabel('PC2')
```

```
ax.set_zlabel('PC3')
```

```
ax.legend()
```

```
plt.show()
```

A 3D Projection of Data In the Reduced Dimension



[Back to Table of Contents](#)


K-means Clustering

Elbow method (clustering)

In cluster analysis, the elbow method is a heuristic used in determining the number of clusters in a data set. The method consists of plotting the explained variation as a function of the number of clusters and picking the elbow of the curve as the number of clusters to use.

```
# Instantiate the clustering model and visualizer
model = KMeans()
visualizer = KElbowVisualizer(model, k=(1,10),size=(1080, 500))

visualizer.fit(pca_df)      # Fit the data to the visualizer
visualizer.show()          # Finalize and render the figure
```



```
<Axes: title={'center': 'Distortion Score Elbow for KMeans
Clustering'}, xlabel='k', ylabel='distortion score'>

kmeans = KMeans(n_clusters = 3, init='k-means++',random_state=42)
kmeans.fit(pca_df)

# Now, print the silhouette score of this model


print("silhouette_score is :",silhouette_score(pca_df, kmeans.labels_,
metric='euclidean'))

silhouette_score is : 0.4033339599157649
```

The Silhouette Score is a commonly used metric to evaluate the quality of clusters in a dataset. It measures how similar each data point is to its own cluster (cohesion) compared to other clusters (separation). The Silhouette Score ranges from -1 to 1, where a higher value indicates better-defined clusters

```
from yellowbrick.cluster import SilhouetteVisualizer
visualizer = SilhouetteVisualizer(kmeans,size=(1080, 500))

visualizer.fit(pca_df)      # Fit the data to the visualizer
visualizer.show()          # Finalize and render the figure
```



```
<Axes: title={'center': 'Silhouette Plot of KMeans Clustering for 1991
Samples in 3 Centers'}, xlabel='silhouette coefficient values',
ylabel='cluster label'>
```

```
kmeans.labels_
```

```
array([0, 2, 2, ..., 2, 2, 0], dtype=int32)
```

```
# Create a 3D scatter plot to visualize the clusters
```

```
fig = plt.figure(figsize=(15, 20))
```

```
ax = fig.add_subplot(111, projection='3d')
```

```
# Scatter plot of data points with color-coded clusters
```

```
ax.scatter(pca_df['PC1_mntwines'], pca_df['PC2_income'],
pca_df['PC3_spent'], c=kmeans.labels_, cmap='rainbow')
```

```
ax.set_xlabel('mntwines')
```


```
ax.set_ylabel('income')
```

```
ax.set_zlabel('spent')
```


```
ax.set_title('K-means Clustering Visualization')
```

```
plt.show()
```


K-means Clustering Visualization



```
for label in np.unique(kmeans.labels_):
    X_ = pca_df[label == kmeans.labels_]
    plt.scatter(X_['PC2_income'], X_['PC3_spent'], label=label)
plt.legend()
plt.xlabel('Income is US($)')
plt.ylabel('Total spent')
plt.show()
```



Main Conclusion from Kmeans


Cluster 0 represents low Income and Low Spent Cluster 1 represents High Income and High Spent Cluster represent Medium level of Income and Medium Spent

Agglomerative Hierarchical Clustering

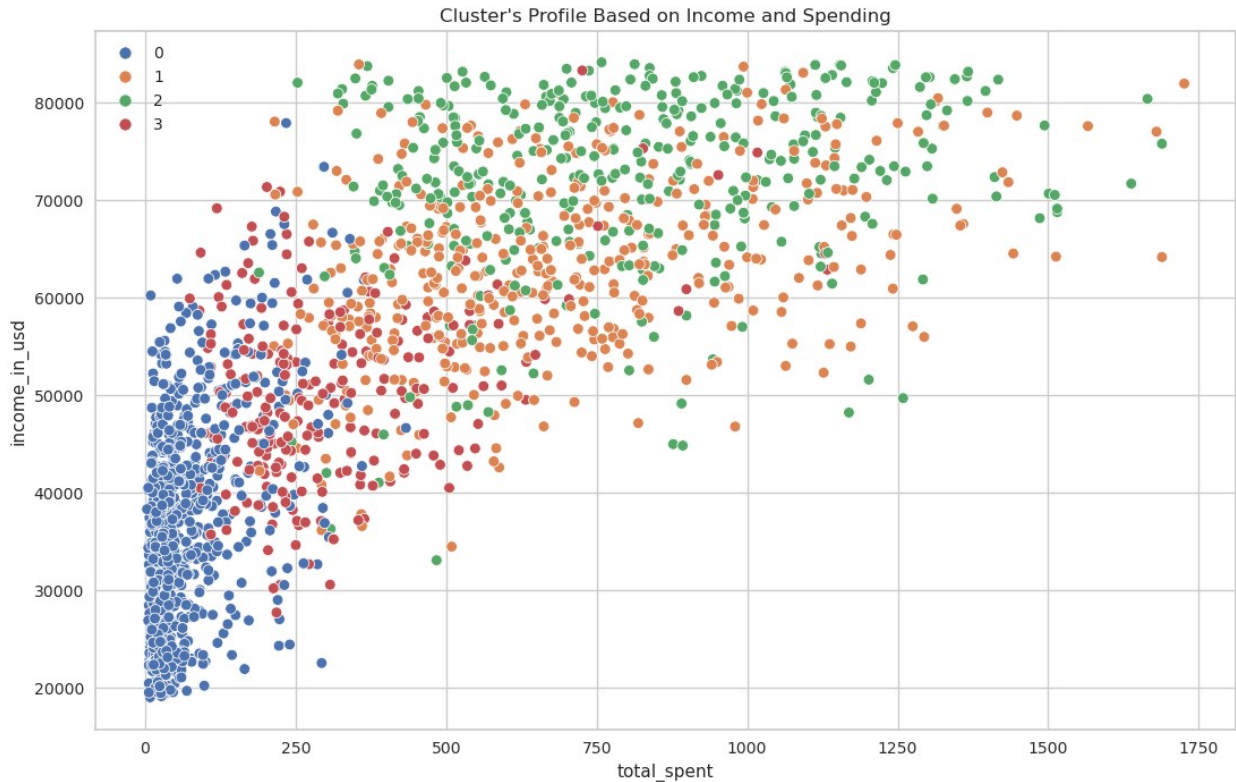
```
AC = AgglomerativeClustering(n_clusters=4)
# fit model and predict clusters
y_hat = AC.fit_predict(pca_df)
pca_df['Clusters'] = y_hat
#Adding the Clusters feature to the original dataframe.
data['Clusters'] = y_hat
old_df['Clusters'] = y_hat

fig = plt.figure(figsize=(13,8))
ax = plt.subplot(111, projection='3d', label='bla')
ax.scatter(x, y, z, s=40, c=pca_df['Clusters'], marker='o',
cmap='Set1_r')
ax.set_title('Clusters')
plt.show()
```

Clusters



```
plt.figure(figsize=(13,8))
pl = sns.scatterplot(data=old_df, x=old_df['total_spent'],
y=old_df['income_in_usd'], hue=old_df['Clusters'], palette= 'deep')
pl.set_title("Cluster's Profile Based on Income and Spending")
plt.legend()
plt.legend();
```



Main Conclusion from Agglomerative Hierarchical Clustering

Cluster0 Low Income Low Spent Cluster1 Mid-High Income Mid-High Spent Cluster2 High income High Spent Cluster3 Low-Mid Income Low-Mid Spent

• DBSCAN

```
cluster = DBSCAN(eps=4, min_samples=4)
cluster.fit(data)
#reducing 1 clustering because its noise/outliers
print(len(set(cluster.labels_) - {1}))
```

2

Finding the % of data marked as noise

```
#-1 is noise
print(f'{100*(cluster.labels_==-1).sum()/len(cluster.labels_)}%')
```

3.2144650929181315%


Visualize the clustered data using matplotlib.pyplot

```
plt.rcParams['figure.figsize'] = (20,15)
unique_labels = set(cluster.labels_)
```

```

n_labels = len(unique_labels)
cmap = plt.cm.get_cmap('brg',n_labels)
for l in unique_labels:
    sns.scatterplot(
        x=data['total_spent'][cluster.labels_==l],
        y=data['income_in_usd'][cluster.labels_==l],
        c = [cmap(l)],
        marker = 'ov'[l%2],
        alpha = 0.75,
        s = 100,
        label = f'Cluster {l} i' if l>=0 else "Noise"
    )
plt.xlabel("Total spent in $")
plt.ylabel("Income in $")
Text(0, 0.5, 'Income in $')

```



Main Conclusion From DBSCAN

We can see only few cluster1 because it is overridden on top by cluster0. However, we can conclude that

- **Cluster 1 represent Mid-High Income and Mid - High Spent**
- **Cluster 0 represents Low - Mid Income and Low Spent**

Among the countries, the top 3 customers with the highest spending are from Spain (48.4%), South Africa (15.9%), and Canada (12.1%). A linear relationship is observed between the 'Total Spent' and 'Income' columns, indicating that as income increases, the total spending also tends to increase. Education level plays a significant role in spending behavior, with postgraduate customers exhibiting the highest spending, followed by graduate and undergraduate customers. Most customers are highly qualified, indicating a potential link between higher education levels and higher spending habits. Utilizing Agglomerative Clustering based on this information yields four distinct clusters: Cluster 0: Low-Income, Low-Spending Cluster 1: Mid-High Income, Mid-High Spending Cluster 2: High Income, High Spending Cluster 3: Low-Mid Income, Low-Mid Spending These findings provide valuable insights for refining marketing strategies, targeting high-value customer segments, and understanding the spending behavior of different income groups.

<h1> Thank You </h1>